

Lecture 14 – Concurrency



Outline:

- ▶ Concurrency vs parallelism
- ▶ Threading
- ▶ Inter-thread communication
- ▶ Asynchrony

Concurrency vs Parallelism



- ▶ Concurrency reduces the latency of a program with IO operations. Concurrent program runs faster than its sequential version even with one physical process.
- ▶ Parallelism reduces the runtime of a program by dividing computation tasks into multiple physical processes.
- ▶ Both concurrency and parallelism use threads and synchronization primitives.
- ▶ Parallel program may be implemented in models such as map-reduce that hides the details of concurrency control.
- ▶ Concurrent programs face a whole host of problems such as space/time leak, deadlock, starvation, and scheduling.

Threading



- ▶ 'forkIO' creates a new thread that runs an IO action. The call returns an IO ThreadId immediately to the parent thread.

```
1 forkIO :: IO () -> IO ThreadID
```

- ▶ Forked child thread races with the parent thread: one creates a file and one tests the existence of a file.

```
1 import Control.Concurrent
2 import System.Directory
3
4     forkIO (writeFile "tmp.txt" "some content")
5 >> doesFileExist "tmp.txt"
```

Reduce latency

- ▶ Perceived latency in interactive applications can be reduced with threads. (Example adopted from Chapter 24 of RWH).

```
1 import qualified Data.ByteString.Lazy as L
2 import Codec.Compression.GZip (compress)
3
4 main = do
5     putStrLn "Enter a file to compress> "
6     line <- getLine           -- get file name
7
8     case line of
9         "" -> return ()        -- treat no name as "want to quit"
10        name ->
11            do content <- L.readFile name
12                -- fork a thread to compress file
13                forkIO (compressFile name content)
14                main           -- run main again without waiting
15 where
16     compressFile path = L.writeFile (path ++ ".gz") . compress
```

Synchronization primitive



- ▶ Just one primitive: MVar – a mutable variable.
"MVar a" can be empty or filed with value a
A thread blocks on a MVar by reading an empty MVar or writing to a full MVar.

```
1 newEmptyMVar :: IO (MVar a) -- creates an empty MVar
2
3 newMVar :: a -> IO (MVar a) -- creates a MVar with value a
4
5 takeMVar :: MVar a -> IO a -- reads from a MVar
6
7 putMVar :: MVar a -> a -> IO () -- writes to a MVar
```

Communicate through MVar



- Below is an example of communication between parent and child thread through MVar.

```
1 communicate = do
2     m <- newEmptyMVar          -- mvar for communication
3
4     forkIO $ do                -- fork a child thread
5         putStrLn ("waiting for parent")
6         v <- takeMVar m        -- block waiting
7         threadDelay 1000000    -- delay child thread 1 second
8         putStrLn ("received " ++ show v)
9
10    threadDelay 1500000       -- delay parent thread 1.5 seconds
11    putStrLn "parent sending"
12    putMVar m "wake up!"      -- send data to MVar
```

- The following output is expected.

```
1 *Main> communicate
2 waiting for parent
3 parent sending
4 (1.50 secs, 24,488 bytes)
5 *Main> received "wake up!"
```

Fetch files 🔊

- ▶ Download and save the first two of the three files

```
1 import Control.Concurrent
2 import qualified Data.ByteString.Lazy as L -- lazy list
3 import Network.HTTP.Conduit           -- http module
4
5 getThree = do
6   v <- newEmptyMVar                  -- mvar to send file to main
7   id1 <- request "http://www.yahoo.com/" v
8   id2 <- request "http://www.google.com/" v
9   id3 <- request "http://www.msn.com/" v
10  takeMVar v >>= L.writeFile "test1.txt"    -- save 1st file
11  takeMVar v >>= L.writeFile "test2.txt"    -- save 2nd file
12  killThread id1                     -- kill them all
13  killThread id2
14  killThread id3
15 where
16   request url v = forkIO $ simpleHttp url -- http request
17                           >>= putMVar v      -- result to v
18                           >> putStrLn url -- print a message
```

Thread functions

- ▶ Create, delay, and destroy a thread.

```
1 -- fork a new thread to run an IO and return thread ID
2 forkIO :: IO () -> IO ThreadId
3
4 -- delay current thread by micro-seconds.
5 threadDelay :: Int -> IO ()
6
7 -- kill thread by its ID
8 -- block until kill signal is received by the target
9 -- returns if the thread is already dead
10 killThread :: ThreadId -> IO ()
```

- ▶ Like MVar functions, every thread function results in an IO action.

Handling Exception 🔊

- ▶ Exception should be handled with catch, handle, or try call.

```
1 import Control.Exception
2
3 -- catch exception 'e' with a handler 'e -> IO a'
4 catch :: Exception e => IO a -> (e -> IO a) -> IO a
5
6 -- 'flip catch'
7 handle :: Exception e => (e -> IO a) -> IO a -> IO a
8
9 -- returns either exception 'e' or result 'a'
10 try :: Exception e => IO a -> IO (Either e a)
11
12 -- turn any exception 'e' into 'SomeException'
13 SomeException :: Exception e => e -> SomeException
14
15 -- 'show' function for exception
16 displayException :: Exception e => e -> String
```

Fetch file with exception handling 🔊

- ▶ Catch exception in forked thread and print its message.

```
1 import Control.Exception
2
3 getThree = do
4     v <- newEmptyMVar                      -- mvar to send file to main
5     id1 <- request "http://www.yahoo.com/"    v
6     id2 <- request "http://www.abcdefghijklmn.com/" v
7     id3 <- request "http://www.msn.com/"      v
8     takeMVar v >>= L.writeFile "test1.txt"   -- save 1st file
9     takeMVar v >>= L.writeFile "test2.txt"   -- save 2nd file
10    killThread id1                         -- kill them all
11    killThread id2
12    killThread id3
13 where
14     request url v = forkIO $(simpleHttp url) -- http request
15                     >>= putMVar v           -- result to v
16                     >>  putStrLn url)  -- print a message
17                     `catch` handler  -- catch exception
18 where handler = \(SomeException e) ->
19     putStrLn $ url ++ ": " ++ (displayException e)
```

- ▶ Can catch both http or thread exception

Fetch file with exception handling 🔊

- ▶ Catch exception in main thread but can't stop the 3rd request

```
1 getThree' = do
2   v <- newEmptyMVar                      -- mvar to send file to main
3   id1 <- request "http://www.yahoo.com/" v -- file or error
4   id2 <- request "http://www.abcccccdffff.com/" v
5   id3 <- request "http://www.msn.com/" v
6   deposit "test1.txt" v                   -- save file or print error
7   deposit "test2.txt" v
8   deposit "test3.txt" v
9   killThread id1                         -- nothing to kill
10  killThread id2
11  killThread id3
12 where
13   deposit file v =
14     takeMVar v >>= either handler (L.writeFile file)
15   request url v = forkIO $
16     try (simpleHttp url) >>= putMVar v >> putStrLn url
17   handler =
18     \(SomeException e) -> putStrLn (displayException e)
```

- ▶ 'try' runs 'IO a' and returns 'either' exception 'e' or result 'a'.

```
1 try :: Exception e => IO a -> IO (Either e a)
```

Concurrency control with MVar

► Exception-safe modification and non-blocking access of MVar

```
1 import Control.Concurrent.MVar
2
3 -- modify a MVar by updating its content 'a'
4 -- and return 'b' (often just to read 'a')
5 modifyMVar :: MVar a -> (a -> IO (a, b)) -> IO b
6
7 -- modify a MVar by updating its content 'a' only
8 modifyMVar_ :: MVar a -> (a -> IO a) -> IO ()
9
10 -- atomically read the content of an MVar
11 readMVar :: MVar a -> IO a
12
13 -- put a new value into the MVar
14 -- and return the value taken
15 swapMVar :: MVar a -> a -> IO a
16
17 -- non-blocking take -- take Nothing or Just a
18 tryTakeMVar :: MVar a -> IO (Maybe a)
19 -- non-blocking put -- True if succeeded in put
20 tryPutMVar :: MVar a -> a -> IO Bool
21 -- non-blocking read -- read Nothing or Just a
22 tryReadMVar :: MVar a -> IO (Maybe a)
```

Concurrency control with MVar

- ▶ Use MVar v1 to count successful requests and MVar v2 to notify main thread to cancel.

```
1 getThree '' = do
2   v1 <- newMVar 0                      -- counter mvar
3   v2 <- newEmptyMVar                   -- cancel mvar
4   id1 <- request "http://www.yahoo.com/" "test1.txt" v1 v2
5   id2 <- request "http://www.google.com/" "test2.txt" v1 v2
6   id3 <- request "http://www.msn.com/" "test3.txt" v1 v2
7
8   takeMVar v2                         -- start cancellation
9   killThread id1; killThread id2; killThread id3
10 where
11   request url f v1 v2 = forkIO $
12     do d <- simpleHttp url      -- update count below
13       count <- modifyMVar v1 (\x -> return (x+1, x+1))
14       putStrLn $ (show count) ++ ":" ++ url
15       L.writeFile f d           -- save the file
16       if count >= 2
17         then putMVar v2 ()    -- send cancel signal
18         else return ()
19   `catch` handler                    -- catch exception
20   where handler = \(SomeException e) ->
21       putStrLn (url ++ ":" ++ displayException e)
```

Communicate with Channels

- ▶ Send/receive messages using Channel, which is an infinite buffer built from MVars.

```
1 import Control.Concurrent.Chan
2
3 -- make a new channel of type 'a'
4 newChan :: IO (Chan a)
5
6 -- write 'a' to the channel and return immediately
7 writeChan :: Chan a -> a -> IO ()
8
9 -- read 'a' from the channel and blocks if empty
10 readChan :: Chan a -> IO a
11
12 -- duplicate a channel, which allows broadcast
13 dupChan :: Chan a -> IO (Chan a)
```

- ▶ Channel is never closed and a thread can block on reading a channel forever even if no thread is sending data to it.

Send messages through channels

- ▶ Send success and exception messages to main via channel.

```
1 getThree''' = do
2     let urls = [("http://www.yahoo.com/", "test1.txt"),
3                  ("http://www.ababbaa.com/", "test2.txt"),
4                  ("http://www.msn.com/", "test3.txt")]
5     v1 <- newMVar 0                      -- counter mvar
6     v2 <- newEmptyMVar                  -- cancel mvar
7     ch <- newChan                      -- message channel
8     ids <- mapM (\(url, path) -> request url path v1 v2 ch) urls
9     takeMVar v2                         -- start cancellation
10    mapM_ killThread ids                -- cancel all
11    mapM_ (\_ -> readChan ch >>= putStrLn) ids
12 where
13     request url f v1 v2 ch = forkIO $
14         do d <- simpleHttp url      -- get file; update counter
15             count <- modifyMVar v1 (\x -> return (x+1, x+1))
16             writeChan ch $ (show count) ++ ":" ++ url
17             L.writeFile f d          -- save file
18             if count >= 2
19                 then putMVar v2 ()  -- signal cancel
20                 else return ()
21             `catch` handler        -- catch exception
22             where handler = \(SomeException e) ->
23                 writeChan ch (url ++ ":" ++ displayException e)
```

Send messages through list of MVars

- ▶ Main thread assigns each child thread an empty MVar.

```
1 getThree''' = do
2     let urls = [("http://www.yahoo.com/", "test1.txt"),
3                  ("http://www.google.com/", "test2.txt"),
4                  ("http://www.msn.com/", "test3.txt")]
5     v1 <- newMVar 0                                -- counter mvar
6     v2 <- newEmptyMVar                            -- cancel mvar
7     vs <- mapM (\_ -> newEmptyMVar) urls        -- message mvars
8     ids <- mapM (\((url, path), v) ->          -- launch threads
9                     request url path v1 v2 v) $ zip urls vs
10    takeMVar v2                                    -- start cancellation
11    mapM killThread ids                          -- cancel all
12    messages <- mapM (tryTakeMVar) vs           -- nonblocking take
13    mapM_ (putStrLn . show) messages            -- print message
14 where
15     request url f v1 v2 v = forkIO $
16         do d <- simpleHttp url
17             count <- modifyMVar v1 (\x -> return (x+1, x+1))
18             putMVar v $ (show count) ++ ":" ++ url
19             L.writeFile f d
20             if count >= 2 then putMVar v2 () else return ()
21             `catch` handler
22             where handler = \(SomeException e) ->
23                   putMVar v $ (url ++ ":" ++ displayException e)
```

Asynchrony



- ▶ Perform asynchronous IO without using threads or MVar.

```
1 import Control.Concurrent.Async
2
3 -- run an action in a separate thread
4 async :: IO a -> IO (Async a)
5
6 -- blocking wait for an async action to return
7 wait :: Async a -> IO a
8
9 -- cancel an async action
10 cancel :: Async a -> IO ()
11
12 -- race two actions and cancel the loser
13 race :: IO a -> IO b -> IO (Either a b)
14 race_ :: IO a -> IO b -> IO ()
15
16 -- run two actions concurrently
17 concurrently :: IO a -> IO b -> IO (a, b)
18 concurrently_ :: IO a -> IO b -> IO ()
19
20 -- run a list of actions concurrently
21 mapConcurrently :: (a -> IO b) -> [a] -> IO [b]
22 mapConcurrently_ :: (a -> IO b) -> [a] -> IO ()
```

Asynchrony



- ▶ Run http requests as a list of async actions

```
1 import Control.Concurrent.Async
2
3 getAsync = do
4     let urls = [("http://www.yahoo.com/", "test1.txt"),
5                 ("http://www.google.com/", "test2.txt"),
6                 ("http://www.abcdefghijklmn.com/", "test3.txt")]
7
8     mapConcurrently_ request urls -- run requests w. no results
9     where
10        request (url, f) = do
11            d <- simpleHttp url
12            putStrLn url
13            L.writeFile f d
14            `catch` handler
15            where handler = \(SomeException e) ->
16                putStrLn (url ++ ": " ++ displayException e)
```

Asynchrony

- ▶ Run http requests as a list of async actions, where messages are printed independently.

```
1 -- output
2 *Main> getAsync
3 http://www.google.com/
4 http://www.abcdefghijklmn.com/: HttpExceptionRequest Request +
5   host                  = "www.abcdefghijklmn.com"
6   port                  = 80
7   secure                = False
8   requestHeaders        = [("Connection", "close")]
9   path                  = "/"
10  queryString          = ""
11  method                = "GET"
12  proxy                 = Nothing
13  rawBody               = False
14  redirectCount         = 10
15  responseTimeout       = ResponseTimeoutDefault
16  requestVersion        = HTTP/1.1
17 }
18 (ConnectionFailure user error (Network.Socket.gai_strerror no
19 http://www.yahoo.com/
20 (1.08 secs, 52,608 bytes)
```

Asynchrony



- ▶ Return messages to the main thread.

```
1 import Control.Concurrent.Async
2
3 getAsync' = do
4     let urls = [("http://www.yahoo.com/", "test1.txt"),
5                 ("http://www.google.com/", "test2.txt"),
6                 ("http://www.abcdefghijklmn.com/", "test3.txt")]
7
8     messages <- mapConcurrently request urls -- run requests
9     mapM_ putStrLn messages                                -- print messages
10    where
11        request (url, f) = do
12            d <- simpleHttp url
13            L.writeFile f d
14            return url
15        `catch` handler
16        where handler = \(SomeException e) ->
17            return (url ++ ":" ++ displayException e)
```

Asynchrony



- ▶ Return messages to the main thread, where messages are printed in the same order as the actions.

```
1 -- output
2 *Main> getAsync'
3 http://www.yahoo.com/
4 http://www.google.com/
5 http://www.abcdefghijklmn.com/: HttpNotFoundException Request +
6   host                      = "www.abcdefghijklmn.com"
7   port                      = 80
8   secure                    = False
9   requestHeaders            = [("Connection", "close")]
10  path                      = "/"
11  queryString               = ""
12  method                    = "GET"
13  proxy                     = Nothing
14  rawBody                   = False
15  redirectCount             = 10
16  responseTimeout           = ResponseTimeoutDefault
17  requestVersion            = HTTP/1.1
18 }
19
20 (ConnectionFailure user error (Network.Socket.gai_strerror no
21 (1.07 secs, 541,480 bytes)
```