

# Lecture 15: Software Transactional Memory

Outline:

- ▶ Review of find file function
- ▶ Software transactional memory (STM)

## Find files in a directory

- ▶ Traverse a file directory and its sub-directories with a state monad, which decides whether the traversal should be *Done*, *Continue*, or *Skip* a sub-directory.

```
1 data FileAction = Done | Skip | Continue
```

- ▶ The result of the traversal is accumulated in State Monad – `FileState s`, where `s` is the state type and contains data that we are interested in.

```
1 type FileState s = StateT s IO FileAction
```

- ▶ The value of `FileState s` is `FileAction`, which tells the `find` function what to do based on the current file and accumulated value (i.e. the state '`s`').

# Find files in a directory



- ▶ The *find* function first lists the file names in the top-level directory of the given *path* and then calls a local function *iterate* to visit each file name.

```
1 find :: forall s.                      -- bind 's' for 'iterate'
2     (FileInfo -> FileState s)  -- getState function
3     -> FilePath                  -- top-level directory
4     -> FileState s              -- final file state
5
6 find getState path = do
7     names <- lift $ listDirectory path --lift IO to StateT s IO
8     iterate names -- recursively visit each file or directory
9
10 where iterate :: [FilePath] -> FileState s
```

- ▶ The *find* function returns a *StateT s IO* monad. So if we want to extract value from an *IO* monad, we have to first 'lift' it using the *lift* function.

```
1 lift :: (MonadTrans t, Monad m) => m s -> t m s
```

# Find files in a directory

- ▶ *iterate* returns *Continue* action when it reaches the end of the list so that *find* function will not stop just because it has searched a sub-directory.

```
1      iterate :: [FilePath] -> FileState s
2
3      iterate [] = return Continue -- continue if finished
4                                         -- searching a directory
5      iterate (name:rest) = do
6          let path' = path </> name           -- full path to file
7          info <- lift $ getInfo path'       -- get file info
8          action <- getState info          -- run file state
9                                         -- to get file action
10
11         case action of
12             Done -> return Done           -- pass Done to caller
13             Skip -> iterate rest        -- don't recurse
14             _ -> if runFileP searchP info -- is a directory
15                 then do -- recursively visit sub-directory
16                     action' <- find getState path'
17                     case action' of
18                         Done -> return Done -- done
19                         _ -> iterate rest -- continue
20                     else iterate rest    --- visit other files
```

# Software Transactional Memory



- ▶ Errors in concurrent programs are often caused by race conditions, where multiple threads have interleaving access to shared variables.
- ▶ To prevent unsafe access to shared variables, concurrent programs use locks; but lock-based programs do not compose, and they may cause deadlock and priority inversion,
- ▶ STM helps manage internal concurrency between the threads interacting through memory in a single process.
  - ▶ A transaction is executed *atomically* – all or nothing.
  - ▶ Only revocable operations (memory access) can be in a *STM*.
  - ▶ Blocking is composable, where a transaction can block at *retry* until it commits.
  - ▶ Can compose alternative transactions using *orElse*.

# Software Transactional Memory 🔊

- ▶ STM is defined as a monad that wraps computation that involves transactional variable TVar.
- ▶ TVars are defined and accessed in STM (except newTVarIO).
- ▶ STM is evaluated *atomically* to IO.

```
1 -- STM monad
2 data STM a
3 instance Monad STM
4
5 -- running STM computations
6 atomically :: STM a -> IO a           -- commits STM a
7 retry      :: STM a          -- retry if anything changes
8 orElse     :: STM a -> STM a -> STM a -- choice of two
9
10 -- transactional variables
11 data TVar a
12 newTVar   :: a -> STM (TVar a)
13 readTVar  :: TVar a -> STM a
14 writeTVar :: TVar a -> a -> STM ()
```

# Software Transactional Memory



- ▶ When a STM is evaluated, the TVars it accessed are logged
  - ▶ when it commits, the logged TVars are checked for atomicity violation (whether any TVar is mutated by another thread).
  - ▶ If no atomicity violation, then STM commits TVar effects and returns a value.
  - ▶ If atomicity violation is detected, all TVar effects are aborted and STM waits for retry (after any TVar is changed).

```
1 -- STM monad
2 data STM a
3 instance Monad STM
4
5 -- running STM computations
6 atomically :: STM a -> IO a
7 retry      :: STM a
8 orElse     :: STM a -> STM a -> STM a
9
10 -- transactional variables
11 data TVar a
12 newTVar   :: a -> STM (TVar a)
13 readTVar  :: TVar a -> STM a
14 writeTVar :: TVar a -> a -> STM ()
```

# Software Transactional Memory

- ▶ STM execute read/write operations on TVars atomically.
- ▶ When invoked atomically, the STM checks that the logged accesses are valid – no concurrent transaction has committed conflicting updates.

```
1 readTVar  :: TVar a -> STM a
2 writeTVar :: TVar a -> a -> STM ()
3 atomically :: STM a -> IO a
4
5 -- read/write resource
6 type Resource = TVar Int
7
8 putR :: Resource -> Int -> STM ()    -- increment r by i
9 putR r i = do v <- readTVar r          -- read value of r
10           writeTVar r (v + i)           -- write r by v + i
11
12 main = do { ...; atomically (putR r 3); ... }
```

# Software Transactional Memory



- ▶ retry aborts a transaction and restarts from the begining after some logged TVar is changed.

```
1 readTVar  :: TVar a -> STM a
2 writeTVar :: TVar a -> a -> STM ()
3 atomically :: STM a -> IO a
4
5 -- read/write resource
6 type Resource = TVar Int
7
8 getR :: Resource -> Int -> STM ()
9 getR r i = do v <- readTVar r
10           if (v < i) then retry    -- abort and retry
11           else writeTVar r (v-i)
12
13 main = do { ...; atomically (getR r 3); ... }
```

# Software Transactional Memory

- ▶ Sequential composition of STMs is an STM.

```
1 -- read/write resource
2 type Resource = TVar Int
3
4 getR :: Resource -> Int -> STM ()
5 getR r i = do v <- readTVar r
6           if (v < i) then retry    -- abort and retry
7           else writeTVar r (v-i)
8
9 main = atomically $ do { getR r1 3; getR r2 7 }
```

- ▶ main blocks if either r1 or r2 lacks resource but there won't a chance of deadlock.

# Software Transactional Memory



- ▶ Composition of alternative STMs is an STM.

```
1 -- read/write resource
2 type Resource = TVar Int
3
4 getR :: Resource -> Int -> STM ()
5 getR r i = do v <- readTVar r
6           if (v < i) then retry    -- abort and retry
7           else writeTVar r (v-i)
8
9 main = atomically $ do { getR r1 3 `orElse` getR r2 7 }
```

- ▶ if *getR r1 3* retries (and aborts), then *getR r2 7* is run; and if the latter also retries, then the entire call retries.

# Software Transactional Memory



- ▶ Use alternative composition to change blocking behavior

```
1 -- read/write resource
2 type Resource = TVar Int
3
4 getR :: Resource -> Int -> STM ()
5 getR r i = do v <- readTVar r
6           if (v < i) then retry      -- abort and retry
7           else writeTVar r (v-i)
8
9 nonblockGetR :: Resource -> Int -> STM Bool
10 nonblockGetR r i = do { getR r i; return True }
11           `orElse` return False
12
13 blockGetR :: Resource -> Int -> STM ()
14 blockGetR r i = do s <- nonBlockGetR r i
15           if s then return () else retry
```

# Software Transactional Memory



- ▶ Composition with `orElse` is associative and has unit `retry`.

```
1 s1 `orElse` (s2 `orElse` s3) = (s1 `orElse` s2) `orElse` s3
2
3 retry `orElse` s = s
4
5 s `orElse` retry = s
```

- ▶ STM is an instance of `MonadPlus`

```
1 instance MonadPlus STM where
2     mplus = orElse
3     mzero = retry
```

# Software Transactional Memory



- ▶ Uncaught exception aborts the transaction with no effect.

```
1 -- Exceptions
2 throwSTM :: Exception e => e -> STM a
3 catchSTM :: Exception e => STM a -> (e -> STM a) -> STM a
4
5 atomically $ do
6   n   <- readTVar v_n
7   lim <- readTVar v_lim
8   writeTVar v_n (n+1)
9
10 if n > lim then throwSTM (AssertionFailed "fail")
11 else if n == lim then retry
12 else return ()
```

# Encode MVar with STM 🔊

- ▶ Operations on MVar can be defined as STM.

```
1 type TMVar a = TVar (Maybe a)
2
3 newEmptyTMVar :: STM (TMVar a)
4 newEmptyTMVar = newTVar Nothing
5
6 takeTMVar :: TMVar a -> STM a
7 takeTMVar mv = do
8     v <- readTVar mv
9     case v of Nothing      -> retry -- block until mv has data
10                Just value -> do writeTVar mv Nothing
11                                return value
```

# Encode MVar with STM



- ▶ Operations on MVar can be defined as STM.

```
1 type TMVar a = TVar (Maybe a)
2
3 putTMVar :: TMVar a -> a -> STM ()
4 putTMVar mv value = do
5   v <- readTVar mv
6   case v of Nothing      -> writeTVar mv $ Just value
7             Just value  -> retry
8
9 tryPutTMVar :: TMVar a -> a -> STM Bool
10 tryPutTMVar mv value = do
11   do { putTMVar mv value; return True }
12   `orElse` return False
```

## Multicast channel

- ▶ Represent buffered data by a linked-list, where each node has a data and a TVar to point to the tail.

```
1 type Chain a = TVar (Item a) -- item pointer
2 type Item a = Empty           -- end item of the list
3             | Full a (Chain a) -- node item of the list
4
5 type MChan a = TVar (Chain a) -- pointer to end item pointer
6 type Port a = TVar (Chain a) -- pointer to front item pointer
7
8 newMChan :: STM (MChan a)
9
10 writeMChan :: MChan a -> a -> STM ()
11
12 newPort :: MChan a -> STM (Port a)
13
14 readPort :: Port a -> STM a
```

- ▶ Write data to the end of the list and read data from the front.

```
1     --> Full a1 --> Full a2 --> Full a3 --> Empty
2     |
3     Port                                MChan
```

# Multicast channel



- ▶ Multiple threads can read from (write into) the same channel.

```
1 type Chain a = TVar (Item a)    -- item pointer
2 type Item a = Empty | Full a (Chain a)
3
4 type MChan a = TVar (Chain a) -- pointer to end item pointer
5 type Port a = TVar (Chain a) -- pointer to front item pointer
6
7 newMChan = do { c <- newTVar Empty; newTVar c }
8
9 writeMChain mc v = do
10   c <- readTVar mc           -- end item pointer
11   c' <- newTVar Empty        -- new end item
12   writeTVar c $ Full v c'   -- insert a node item
13   writeTVar mc c'          -- update channel
```

# Multicast channel



- ▶ Multiple threads can read from (write into) the same channel.

```
1 type Chain a = TVar (Item a) -- item pointer
2 type Item a = Empty | Full a (Chain a)
3
4 type MChan a = TVar (Chain a) -- pointer to end item pointer
5 type Port a = TVar (Chain a) -- pointer to front item pointer
6
7 newPort mc = do { c <- newTVar mc; newTVar c }
8
9 readPort p = do
10    c <- readTVar p           -- front item pointer
11    i <- readTVar c           -- front item
12    case i of Empty -> retry -- block on empty
13        Full v c' ->
14            do writeTVar p c' -- update port
15            return v          -- return data
```

# Multicast channel



## ► Merging operations on channel

```
1 type Chain a = TVar (Item a)
2 type Item a = Empty | Full a (Chain a)
3 type MChan a = TVar (Chain a)
4 type Port a = TVar (Chain a)
5
6 -- read from channel p1 or channel p2
7 atomically $ readPort p1 `orElse` readPort p2
8
9 -- read from channel p1 or terminate on signal from m1
10 atomically $ readPort p1 `orElse` takeTMVar m1
11
12 -- merge values from a list of STMs
13 merge :: [STM a] -> STM a
14 merge = foldr1 orElse
```

## Top-level TVar

- ▶ Top-level TVars created with *newTVarIO* allow explicit sharing of TVars between threads.

```
1 newTVar :: a -> STM (TVar a)    -- TVar used in STM
2
3 newTVarIO :: a -> IO (TVar a)   -- TVar used in IO
4
5 readTVar :: TVar a -> STM a     -- read TVar in STM
6
7 readTVarIO :: TVar a -> IO a    -- read TVar in IO
8 -- equal to atomically . readTVar but faster
```